



# CSC415: Introduction to Reinforcement Learning

## Lecture 2: Planning by Dynamic Programming

Dr. Amey Pore

Winter 2026

January 14, 2026

## Think Pair Wise 1

$s_1$	$s_2$	$s_3$	$s_4$	$s_5$	$s_6$	$s_7$
						

In a Markov decision process, a large discount factor  $\gamma$  means that short term rewards are much more influential than long term rewards.

**Answers:** • True   • False   • Don't know

# Outline

- ① Recap
- ② MDPs
- ③ Introduction to Dynamic Programming
- ④ Policy Iteration
- ⑤ Course logistics
- ⑥ Value Iteration

# Return & Value Function

- **Definition of Horizon ( $H$ )**

- Number of time steps in each episode
- Can be infinite
- Otherwise called **finite** Markov reward process

- **Definition of Return,  $G_t$  (for a MRP)**

- Discounted sum of rewards from time step  $t$  to horizon  $H$

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{H-1} r_{t+H-1}$$

- **Definition of State Value Function,  $V(s)$  (for a MRP)**

- Expected return from starting in state  $s$

$$V(s) = \mathbb{E}[G_t | S_t = s] = \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \cdots + \gamma^{H-1} r_{t+H-1} | S_t = s]$$

# Bellman Equation for MRPs

The value function can be decomposed into two parts:

- immediate reward  $R_{t+1}$
- discounted value of successor state  $\gamma V(S_{t+1})$

$$\begin{aligned}V(s) &= \mathbb{E}[G_t | S_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} + \dots) | S_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\&= \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]\end{aligned}$$

## Bellman Equation for MRPs (2)

$$V(s) = \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$$

[Diagram showing state transitions:  $s \rightarrow s'$  with reward  $r$ ]

$$V(s) = \underbrace{R(s)}_{\text{Immediate reward}} + \underbrace{\gamma \sum_{s' \in S} P(s'|s)V(s')}_{\text{Discounted sum of future rewards}}$$

## Matrix Form of Bellman Equation for MRP

For finite state MRP, we can express  $V(s)$  using a matrix equation

$$\begin{pmatrix} V(s_1) \\ \vdots \\ V(s_N) \end{pmatrix} = \begin{pmatrix} R(s_1) \\ \vdots \\ R(s_N) \end{pmatrix} + \gamma \begin{pmatrix} P(s_1|s_1) & \cdots & P(s_N|s_1) \\ P(s_1|s_2) & \cdots & P(s_N|s_2) \\ \vdots & \ddots & \vdots \\ P(s_1|s_N) & \cdots & P(s_N|s_N) \end{pmatrix} \begin{pmatrix} V(s_1) \\ \vdots \\ V(s_N) \end{pmatrix}$$

$$V = R + \gamma PV$$

## Analytic Solution for Value of MRP

For finite state MRP, we can express  $V(s)$  using a matrix equation

$$V = R + \gamma PV$$

$$V - \gamma PV = R$$

$$(I - \gamma P)V = R$$

$$V = (I - \gamma P)^{-1}R$$

- Solving directly requires taking a matrix inverse  $\sim O(N^3)$
- Requires that  $(I - \gamma P)$  is invertible
- Direct solutions only possible for small MRPs

# Markov Decision Process (MDP)

A Markov decision process (MDP) is a Markov reward process with decisions. It is an *environment* in which all states are Markov.

## Definition

A Markov Decision Process is a tuple  $(\mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma)$

- $\mathcal{S}$  is a (finite) set of Markov states  $s \in \mathcal{S}$
- $\mathcal{A}$  is a (finite) set of actions  $a \in \mathcal{A}$
- $\mathcal{P}$  is dynamics/transition model for **each action**,  $P(s_{t+1} = s' | s_t = s, a_t = a)$
- $\mathcal{R}$  is a reward function  $R(s_t = s, a_t = a) = \mathbb{E}[r_t | s_t = s, a_t = a]$
- Discount factor  $\gamma \in [0, 1]$

# Example: Mars Rover MDP

$s_1$	$s_2$	$s_3$	$s_4$	$s_5$	$s_6$	$s_7$
						

- 2 deterministic actions

$$P(s'|s, a_1) = \begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}$$

$$P(s'|s, a_2) = \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$$

# MDP Policies

## Definition

A *policy*  $\pi$  is a distribution over actions given states,

$$\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$$

- A policy fully defines the behaviour of an agent
- MDP policies depend on the current state (not the history)
- i.e. Policies are *stationary* (time-independent),

$$A_t \sim \pi(\cdot | S_t), \forall t > 0$$

## MDP Policies (2)

- Given an MDP  $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and a policy  $\pi$
- The state sequence  $S_1, S_2, \dots$  is a Markov process  $\langle \mathcal{S}, \mathcal{P}^\pi \rangle$
- The state and reward sequence  $S_1, R_2, S_2, \dots$  is a Markov reward process  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$
- where

$$R^\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) R(s, a)$$

$$P^\pi(s'|s) = \sum_{a \in \mathcal{A}} \pi(a|s) P(s'|s, a)$$

# Value Function under a Policy


## Definition

The state-value function  $V^\pi(s)$  of an MDP is the expected return starting from state  $s$ , and then following policy  $\pi$

$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s]$$

# Bellman Expectation Equation

The state-value function can again be decomposed into immediate reward plus discounted value of successor state,  $V^\pi(s) = \mathbb{E}_\pi[r_{t+1} + \gamma V^\pi(s_{t+1})|s_t = s]$



$$V^\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left[ R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^\pi(s') \right]$$

# Optimal Value Function

## Definition

The *optimal state-value function*  $V^*(s)$  is the maximum value function over all policies

$$V^*(s) = \max_{\pi} V^{\pi}(s)$$

- The optimal value function gives the best possible performance in the MDP
- An MDP is "solved" when we know the optimal value function

# Optimal Policy

Define a partial ordering over policies

$$\pi \geq \pi' \text{ if } V^\pi(s) \geq V^{\pi'}(s), \forall s$$

## Theorem

For any Markov Decision Process

- There exists an optimal policy  $\pi^*$  that is better than or equal to all other policies,  
 $\pi^* \geq \pi, \forall \pi$
- All optimal policies achieve the optimal value function,

$$V^{\pi^*}(s) = V^*(s)$$

# Finding an Optimal Policy

- Compute the optimal policy

$$\pi^*(s) = \arg \max_{\pi} V^{\pi}(s)$$

- There exists a **unique** optimal value function
- Optimal policy for a MDP in an infinite horizon problem is deterministic


# Finding an Optimal Policy

- Compute the optimal policy

$$\pi^*(s) = \arg \max_{\pi} V^{\pi}(s)$$

- There exists a **unique** optimal value function
- Optimal policy for a MDP in an infinite horizon problem (agent acts forever) is:
  - Deterministic
  - Stationary (does not depend on time step)
  - Unique? Not necessarily, may have two policies with identical (optimal) values

# Bellman Optimality Equation



$$V^*(s) = \max_a R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s')$$

# Today's lecture: Recommendations

- The content is theoretical, but fundamental to RL
- Hence, if you don't ask questions, you will not understand
- You might get a reward for asking questions

# What is Dynamic Programming?

- **Dynamic:** sequential or temporal component to the problem
- **Programming:** optimising a “program”, i.e. a policy
  - c.f. linear programming
- A method for solving complex problems
- By breaking them down into subproblems
  - Solve the subproblems
  - Combine solutions to subproblems

# Requirements for Dynamic Programming

Dynamic Programming is a very general solution method for problems which have two properties:

- **Optimal substructure**

- Principle of optimality applies
- Optimal solution can be decomposed into subproblems

- **Overlapping subproblems**

- Subproblems recur many times
- Solutions can be cached and reused

- **MDPs satisfy both properties**

- Bellman equation gives recursive decomposition
- Value function stores and reuses solutions


# Planning by Dynamic Programming

- Dynamic programming assumes **full knowledge of the MDP**
- It is used for **planning** in an MDP
- **For prediction:**
  - Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$  and policy  $\pi$
  - or: MRP  $\langle \mathcal{S}, \mathcal{P}^\pi, \mathcal{R}^\pi, \gamma \rangle$
  - Output: value function  $V^\pi$
- **Or for control:**
  - Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
  - Output: optimal value function  $V^*$
  - and: optimal policy  $\pi^*$

# Iterative Policy Evaluation


- **Problem:** evaluate a given policy  $\pi$
- **Solution:** iterative application of Bellman expectation backup
- $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_\pi$
- Using **synchronous** backups,
  - At each iteration  $k + 1$
  - For all states  $s \in \mathcal{S}$
  - Update  $V_{k+1}(s)$  from  $V_k(s')$
  - where  $s'$  is a successor state of  $s$

# Iterative Policy Evaluation 2



$$V_{k+1}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_k(s') \right)$$

# Evaluating a Random Policy in the Small Gridworld




	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$r = -1$   
on all transitions

- Undiscounted episodic MDP ( $\gamma = 1$ )
- Nonterminal states  $1, \dots, 14$
- One terminal state (shown twice as shaded squares)
- Actions leading out of the grid leave state unchanged
- Reward is  $-1$  until the terminal state is reached
- Agent follows uniform random policy

$$\pi(n|\cdot) = \pi(e|\cdot) = \pi(s|\cdot) = \pi(w|\cdot) = 0.25$$

# Iterative Policy Evaluation in Small Gridworld



# Iterative Policy Evaluation in Small Gridworld (2)

 $k = 3$ 




0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

 $k = 10$ 

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

 $k = \infty$ 

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0



optimal  
policy

## Homework Exercise: MDP 1 Iteration of Policy Evaluation

- **Dynamics:**  $P(s_6|s_6, a_1) = 0.5, P(s_7|s_6, a_1) = 0.5, \dots$
- **Reward:** for all actions, +1 in state  $s_1$ , +10 in state  $s_7$ , 0 otherwise
- Let  $\pi(s) = a_1 \forall s$ , assume  $V_k = [1, 0, 0, 0, 0, 0, 10]$  and  $k = 1, \gamma = 0.5$
- **Compute**  $V_{k+1}(s_6)$

*See answer at the end of the slide deck. Work this out yourself, if you'd like to practice. Then check the answer.*

# How to Improve a Policy

- Given a policy  $\pi$

- Evaluate the policy  $\pi$


$$V^\pi(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$$

- Improve the policy by acting greedily with respect to  $V^\pi$


$$\pi' = \text{greedy}(V^\pi)$$

- In Small Gridworld improved policy was optimal,  $\pi' = \pi^*$
- In general, need more iterations of improvement / evaluation
- But this process of **policy iteration** always converges to  $\pi^*$

# Policy Iteration



- Policy evaluation Estimate  $V^\pi$   
Iterative policy evaluation
  - Policy improvement Generate  $\pi' \geq \pi$   
Greedy policy improvement



# New Definition: State-Action Value Q

- State-action value of a policy

$$Q^\pi(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^\pi(s')$$

- Take action  $a$ , then follow the policy  $\pi$

# Policy Improvement

- Compute state-action value of a policy  $\pi_i$ :
  - For  $s$  in  $S$  and  $a$  in  $A$ :

$$Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^{\pi_i}(s')$$

- Compute new policy  $\pi_{i+1}$ , for all  $s \in S$

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a) \quad \forall s \in S$$

# MDP Policy Iteration (PI)

- Set  $i = 0$
- Initialize  $\pi_0(s)$  randomly for all states  $s$
- While  $i == 0$  or  $\|\pi_i - \pi_{i-1}\|_1 > 0$  (L1-norm, measures if the policy changed for any state):
  - $V^{\pi_i} \leftarrow \text{MDP } V \text{ function policy evaluation of } \pi_i$
  - $\pi_{i+1} \leftarrow \text{Policy improvement}$
  - $i = i + 1$

# Delving Deeper Into Policy Improvement Step

$$Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^{\pi_i}(s')$$

# Delving Deeper Into Policy Improvement Step

$$Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^{\pi_i}(s')$$

$$\max_a Q^{\pi_i}(s, a) \geq R(s, \pi_i(s)) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, \pi_i(s)) V^{\pi_i}(s') = V^{\pi_i}(s)$$

$$\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$$

- Suppose we take  $\pi_{i+1}(s)$  for one action, then follow  $\pi_i$  forever
  - Our expected sum of rewards is at least as good as if we had always followed  $\pi_i$
- But new proposed policy is to always follow  $\pi_{i+1} \dots$

# Monotonic Improvement in Policy

## Definition

$$V^{\pi_1} \geq V^{\pi_2}: V^{\pi_1}(s) \geq V^{\pi_2}(s), \forall s \in \mathcal{S}$$

**Proposition:**  $V^{\pi_{i+1}} \geq V^{\pi_i}$  with strict inequality if  $\pi_i$  is suboptimal, where  $\pi_{i+1}$  is the new policy we get from policy improvement on  $\pi_i$

# Proof: Monotonic Improvement in Policy

$$\begin{aligned} V^{\pi_i}(s) &\leq Q^{\pi_i}(s, \pi_{i+1}(s)) \\ &= \mathbb{E}_{\pi_{i+1}}[R_{t+1} + \gamma V^{\pi_i}(S_{t+1}) | S_t = s] \\ &\leq \mathbb{E}_{\pi_{i+1}}[R_{t+1} + \gamma Q^{\pi_i}(S_{t+1}, \pi_{i+1}(S_{t+1})) | S_t = s] \\ &= \mathbb{E}_{\pi_{i+1}}[R_{t+1} + \gamma R_{t+2} + \gamma^2 V^{\pi_i}(S_{t+2}) | S_t = s] \\ &\leq \mathbb{E}_{\pi_{i+1}}[R_{t+1} + \gamma R_{t+2} + \dots] \\ &= V^{\pi_{i+1}}(s) \end{aligned}$$

# Course logistics

- ➊ First lab tomorrow: dynamic programming, due date on Jan 20th at 11:59pm.
  - The submission will be on MarkUs.
  - Submit .ipynb notebook.
- ➋ The office hours time for the TA has changed to Tuesday 5:15pm to 6:15pm.
  - Link to the zoom will be added on the home page Quercus.
- ➌ Mid-term exam will take place on Jan 29th covering the first 4 lectures.
  - **Important:** We will need to exceed the tutorial time by 30mins. It is a 90 mins exam.
  - I will add sample questions on Piazza later today.

## Think Pair Wise 2

- **If policy doesn't change, can it ever change again?**
  - Yes / No / Not sure
- **Is there a maximum number of iterations of policy iteration?**
  - Yes / No / Not sure

# Break!

# Outline

- ① Recap
- ② Introduction to Dynamic Programming
- ③ Policy Iteration
- ④ **Value Iteration**

# Principle of Optimality

Any optimal policy can be subdivided into two components:

- An optimal first action  $A_*$
- Followed by an optimal policy from successor state  $S'$

## Theorem (Principle of Optimality)

A policy  $\pi(a|s)$  achieves the optimal value from state  $s$ ,  $V^\pi(s) = V^*(s)$ , if and only if

- For any state  $s'$  reachable from  $s$
- $\pi$  achieves the optimal value from state  $s'$ ,  $V^\pi(s') = V^*(s')$

# Deterministic Value Iteration

- If we know the solution to subproblems  $V^*(s')$
- Then solution  $V^*(s)$  can be found by one-step lookahead

$$V^*(s) \leftarrow \max_{a \in \mathcal{A}} \left[ R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^*(s') \right]$$

- The idea of value iteration is to apply these updates iteratively
- Intuition: start with final rewards and work backwards
- Still works with loopy, stochastic MDPs

# Example: Shortest Path

$g$				

Problem

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

 $V_1$ 

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1

 $V_2$ 

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-2
-2	-2	-2	-2

 $V_3$ 

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3

 $V_4$ 

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-4
-3	-4	-4	-4

 $V_5$ 

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-5

 $V_6$ 


0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

 $V_7$

# Value Iteration

- Problem: find optimal policy  $\pi$
- Solution: iterative application of Bellman optimality backup
- $V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V^*$
- Using synchronous backups
  - At each iteration  $k + 1$
  - For all states  $s \in \mathcal{S}$
  - Update  $V_{k+1}(s)$  from  $V_k(s')$
- Convergence to  $V^*$  will be proven later
- Unlike policy iteration, there is no explicit policy
- Intermediate value functions may not correspond to any policy

## Value Iteration (2)



$$V_{k+1}(s) = \max_{a \in \mathcal{A}} \left[ R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V_k(s') \right]$$

# Example of Value Iteration in Practice

<https://artint.info/demos/mdp/vi.html>

# Going Back to Value Iteration (VI)

- Set  $k = 1$
- Initialize  $V_0(s) = 0$  for all states  $s$
- Loop until convergence: (for ex.  $\|V_{k+1} - V_k\|_\infty \leq \epsilon$ )
  - For each state  $s$

$$V_{k+1}(s) = \max_a \left[ R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_k(s') \right]$$

- To extract optimal policy if can act for  $k + 1$  more steps,

$$\pi(s) = \arg \max_a \left[ R(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_{k+1}(s') \right]$$

# Synchronous Dynamic Programming Algorithms

Problem	Bellman Equation	Algorithm
Prediction	Bellman Expectation Equation	Iterative Policy Evaluation
Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
Control	Bellman Optimality Equation	Value Iteration

- Algorithms are based on state-value function  $V^\pi(s)$  or  $V^*(s)$
- Complexity  $O(mn^2)$  per iteration, for  $m$  actions and  $n$  states
- Could also apply to action-value function  $Q^\pi(s, a)$  or  $Q^*(s, a)$
- Complexity  $O(m^2n^2)$  per iteration

# Value vs Policy Iteration

	Policy Iteration	Value Iteration
Algorithm	Iterates through evaluation and improvement	Iterates through Bellman optimality updates
Updates	Solves $V^\pi$ for a fixed policy	Updates $V(s)$ using max over a
Convergence	Fewer iterations, but each is costly	More iterations, but each is simple
Output	Directly improves the policy	Extracts policy from optimal values

# What You Should Know

- Define MP, MRP, MDP, Bellman equation, model, Q-value, policy
- Be able to implement
  - Value Iteration
  - Policy Iteration
- Give pros and cons of different policy evaluation approaches
- Be able to prove contraction properties
- Limitations of presented approaches and Markov assumptions

# Thank you!

# Some Technical Questions

- How do we know that value iteration converges to  $V^*$ ?
- Or that iterative policy evaluation converges to  $V^\pi$ ?
- And therefore that policy iteration converges to  $V^*$ ?
- Is the solution unique?
- How fast do these algorithms converge?
- These questions are resolved by *contraction mapping theorem*

# Value Function Space

- Consider the vector space  $\mathcal{V}$  over value functions
- There are  $|\mathcal{S}|$  dimensions
- Each point in this space fully specifies a value function  $V(s)$
- What does a Bellman backup do to points in this space?
- We will show that it brings value functions *closer*
- And therefore the backups must converge on a unique solution

# Value Function $\infty$ -Norm

- We will measure distance between state-value functions  $U$  and  $V$  by the  $\infty$ -norm
- i.e. the largest difference between state values,

$$\|U - V\|_{\infty} = \max_{s \in \mathcal{S}} |U(s) - V(s)|$$

# Bellman Expectation Backup is a Contraction

- Define the *Bellman expectation backup operator*  $T^\pi$ ,

$$T^\pi(V) = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi V$$

# Bellman Expectation Backup is a Contraction

- Define the *Bellman expectation backup operator*  $T^\pi$ ,

$$T^\pi(V) = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi V$$

- This operator is a  $\gamma$ -contraction, i.e. it makes value functions closer by at least  $\gamma$ ,

$$\begin{aligned}\|T^\pi(U) - T^\pi(V)\|_\infty &= \|(\mathcal{R}^\pi + \gamma \mathcal{P}^\pi U) - (\mathcal{R}^\pi + \gamma \mathcal{P}^\pi V)\|_\infty \\ &= \|\gamma \mathcal{P}^\pi(U - V)\|_\infty \\ &\leq \|\gamma \mathcal{P}^\pi\|_\infty \|U - V\|_\infty \\ &\leq \gamma \|U - V\|_\infty\end{aligned}$$

# Contraction Mapping Theorem

## Theorem (Contraction Mapping Theorem)

For any metric space  $\mathcal{V}$  that is complete (i.e. closed) under an operator  $T(\mathcal{V})$ , where  $T$  is a  $\gamma$ -contraction,

- $T$  converges to a unique fixed point
- At a linear convergence rate of  $\gamma$

# Convergence of Iter. Policy Evaluation and Policy Iteration

- The Bellman expectation operator  $T^\pi$  has a unique fixed point
- $V^\pi$  is a fixed point of  $T^\pi$  (by Bellman expectation equation)
- By contraction mapping theorem
- Iterative policy evaluation converges on  $V^\pi$
- Policy iteration converges on  $V^*$

# Bellman Optimality Backup is a Contraction

- Define the *Bellman optimality backup operator*  $T^*$ ,

$$T^*(V) = \max_{a \in \mathcal{A}} \mathcal{R}^a + \gamma \mathcal{P}^a V$$

- This operator is a  $\gamma$ -contraction, i.e. it makes value functions closer by at least  $\gamma$  (similar to previous proof)

$$\|T^*(U) - T^*(V)\|_\infty \leq \gamma \|U - V\|_\infty$$

# Convergence of Value Iteration

- The Bellman optimality operator  $T^*$  has a unique fixed point
- $V^*$  is a fixed point of  $T^*$  (by Bellman optimality equation)
- By contraction mapping theorem
- Value iteration converges on  $V^*$

# Homework Solution

- **Dynamics:**  $P(s_6|s_6, a_1) = 0.5, P(s_7|s_6, a_1) = 0.5, \dots$
- **Reward:** for all actions, +1 in state  $s_1$ , +10 in state  $s_7$ , 0 otherwise
- Let  $\pi(s) = a_1 \forall s$ , assume  $V_k = [1, 0, 0, 0, 0, 0, 10]$  and  $k = 1, \gamma = 0.5$
- **Compute**  $V_{k+1}(s_6)$

$$\begin{aligned}V_{k+1}(s_6) &= R(s_6) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s_6, a_1) V_k(s') \\&= 0 + 0.5 \times (0.5 \times 10 + 0.5 \times 0) \\&= 2.5\end{aligned}$$

## Think pair wise 2: Explanation of Policy Not Changing

- Suppose for all  $s \in \mathcal{S}$ ,  $\pi_{i+1}(s) = \pi_i(s)$
- Then for all  $s \in \mathcal{S}$ ,  $Q^{\pi_{i+1}}(s, a) = Q^{\pi_i}(s, a)$

Recall policy improvement step:

- $Q^{\pi_i}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P(s'|s, a) V^{\pi_i}(s')$
- $\pi_{i+1}(s) = \arg \max_a Q^{\pi_i}(s, a)$
- $\pi_{i+2}(s) = \arg \max_a Q^{\pi_{i+1}}(s, a) = \arg \max_a Q^{\pi_i}(s, a)$

# Opportunities for Out-of-Class Practice

- Does the initialization of values in value iteration impact anything?
- Is the value of the policy extracted from value iteration at each round guaranteed to monotonically improve (if executed in the real infinite horizon problem), like policy iteration?